

VENKATA SANDEEP KUMAR REDDY

DATA ENGINEER | CLOUD DATA ENGINEERING | STREAMING | ANALYTICS ENGINEERING

India | +91 7815930790 | cvsandeep0105@gmail.com | LinkedIn | GitHub | Portfolio

PROFESSIONAL SUMMARY

Data Engineer with 2+ years of professional experience in SQL, data extraction, transformation, validation and ETL workflows, plus hands-on engineering across real-time streaming, cloud-oriented data platforms, CDC, lakehouse architecture, modern data warehousing and analytics engineering. Built three end-to-end data platforms using Kafka, Spark, PostgreSQL, Apache Iceberg, dimensional modeling, ELT, data quality, governance, observability and CI/CD.

CORE TECHNICAL SKILLS

Programming: Python, SQL, PySpark | Data Engineering: ETL, ELT, Batch, Streaming, CDC, Data Quality, Validation, Reconciliation, Data Modeling, Lineage

Processing: Apache Kafka, Apache Spark, Spark Streaming | Platforms: PostgreSQL, Parquet, Apache Iceberg, Data Lakehouse, Data Warehouse

Cloud: AWS-oriented Data Engineering, S3, Glue Data Catalog, Athena, IAM, KMS, CloudWatch, Terraform | Analytics: Dimensional Modeling, SCD, Incremental Models, Data Marts, Semantic Layer

Engineering: Docker, Kubernetes, CI/CD, GitHub Actions, Testing, Observability | Applications: FastAPI, React, Streamlit, scikit-learn

PROFESSIONAL EXPERIENCE

Junior Data Engineer | Mar Symantech Private Limited

December 2023 – August 2026

- Developed and supported data extraction, transformation and ETL workflows using SQL and Python, working with structured datasets and databases to prepare reliable data for downstream reporting and analytics.
- Built and maintained batch and distributed data-processing workflows using Apache Spark/PySpark, applying joins, filtering, aggregation, transformation and source-to-target validation across pipeline stages.
- Worked with Kafka-based data ingestion and streaming workflows, supporting the movement and processing of structured data alongside batch-oriented pipelines.
- Used AWS data services, Amazon S3, Airflow, Docker, Kubernetes and CI/CD practices to support data-processing workflows, orchestration, development, testing and deployment activities.
- Implemented data-quality checks, pipeline validation and troubleshooting, investigating failed or inconsistent processing results and collaborating with senior team members to resolve issues and improve workflow reliability.
- Contributed to development, testing, documentation and pipeline support, following team-defined engineering practices and validating data before downstream consumption.

ENGINEERING PROJECTS

Enterprise AI-Powered Real-Time EV Fleet Data Platform

Real-Time Data Engineering | Streaming | AI-enabled Data Systems

- Built an end-to-end EV telemetry platform using Python, Kafka, Spark, PostgreSQL, FastAPI, Docker and Kubernetes.
- Processed 98,061,214 telemetry records: 98,060,402 valid canonical records and 812 rejected records; generated 160,774 feature windows and 8,251 AI samples.
- Integrated Isolation Forest with 0.8619 F1 and 1.0000 recall; validated 20/20 API benchmark requests at 77.278 ms p95 latency.

Cloud-Native Lakehouse, CDC & AI-Ready Data Platform

Cloud Data Engineering | Lakehouse | CDC

- Engineered a cloud-oriented platform spanning PostgreSQL, raw storage, Bronze Parquet, Apache Iceberg, Silver processing and Gold analytical products.
- Implemented and validated PostgreSQL logical CDC using pgoutput with INSERT/UPDATE/DELETE capture, ordering, LSN progression, deduplication, idempotency and checkpoints.
- Validated 15/15 automated engineering tests, 63/63 data-quality checks and 60/60 reconciliation checks; produced 6 Gold, 9 Raw and 9 Bronze datasets.

RetailIQ — Modern Data Warehouse & Analytics Engineering Platform

Analytics Engineering | ELT | Data Warehousing | BI

- Built a PostgreSQL modern data warehouse from raw/staging layers through dimensional modeling, fact/dimension tables, SCD, incremental models, data marts and semantic views.
- Implemented ELT and analytics-engineering workflows for reusable transformations, analytical business logic and BI-ready products.
- Validated 34/34 data-quality checks and 9/9 integration tests with E2E validation passing; built a six-page Streamlit analytics dashboard and implemented lineage, governance, observability and CI/CD.

EDUCATION

Bachelor of Technology (B.Tech) — Data Science | Presidency University | 2019 – 2023

CERTIFICATIONS & TRAINING

PGP in Data Science and Machine Learning — Intellipaat, 2024 | Linux Certification Training — Intellipaat, 2024 | Data Science Course — Intellipaat, 2023

PROFESSIONAL LINKS

LinkedIn: <https://www.linkedin.com/in/sandeep-reddy-1a9092386>

GitHub: <https://github.com/cvsandeep0105-netizen>

Portfolio: <https://cvsandeep0105-netizen.github.io/data-engineering-portfolio/>