VENKATA SANDEEP KUMAR REDDY

DATA ENGINEER | CLOUD DATA ENGINEERING | STREAMING | ANALYTICS ENGINEERING

India • +91 7815930790 • cvsandeep0105@gmail.com
LinkedIn • GitHub • Portfolio

PROFESSIONAL SUMMARY

Data Engineer with 3 years of professional experience in SQL/Python data workflows and a portfolio spanning real-time streaming, CDC/lakehouse engineering, and analytics engineering. Focus areas include ETL/ELT, distributed processing, data quality, reconciliation, dimensional modeling, testing, and production-oriented platform practices.

TECHNICAL SKILLS

Programming: Python, SQL, PySpark **Streaming/Compute:** Apache Kafka, Spark, Spark Streaming

Data Engineering: ETL/ELT, CDC, incremental processing, deduplication, idempotency, replay, data contracts, quality, reconciliation, lineage

Storage/Modeling: PostgreSQL, Parquet, Apache Iceberg, lakehouse/warehouse, dimensional modeling, facts/dimensions, SCD, data marts, semantic layer

Cloud/Engineering: AWS S3, Glue Data Catalog, Athena, IAM, KMS, CloudWatch, Terraform, Docker, Kubernetes, GitHub Actions, CI/CD, Airflow

APIs/Analytics: FastAPI, React, Streamlit, scikit-learn, Isolation Forest, automated testing, observability, performance benchmarking

PROFESSIONAL EXPERIENCE

Junior Data Engineer | Mar Symantech Private Limited | September 2023 – September 2026

- Developed and supported SQL/Python workflows for extraction, transformation, validation, and downstream analytics. • Built batch/distributed processing with Spark/PySpark and source-to-target checks.
- Worked with Kafka ingestion and cloud-oriented data workflows; used AWS S3, Airflow, Docker, Kubernetes, and CI/CD practices in development/testing workflows.
- Implemented data-quality checks, investigated inconsistent results, and supported troubleshooting, documentation, and workflow reliability.

SELECTED ENGINEERING PROJECTS

Enterprise AI-Powered Real-Time EV Fleet Data Platform

Python, Kafka, Spark, PostgreSQL, FastAPI, Docker, Kubernetes. Processed **98,061,214** telemetry records (**98,060,402 valid**; 812 rejected), producing 160,774 feature windows and 8,251 AI samples. Isolation Forest F1: **0.8619**; API benchmark: **20/20** successful requests at **77.278 ms p95**.

Cloud-Native Lakehouse, CDC & AI-Ready Data Platform

PostgreSQL pgoutput CDC for INSERT/UPDATE/DELETE, LSN progression, ordering, deduplication, idempotency, and checkpoints across Raw/Bronze/Silver/Gold. Recorded evidence: **15/15** automated tests, **63/63** data-quality checks, **60/60** reconciliation checks, 9 Raw datasets, 9 Bronze datasets, and 6 Gold products.

RetailIQ — Modern Data Warehouse & Analytics Engineering Platform

PostgreSQL ELT, dimensional facts/dimensions, SCD, incremental models, data marts, semantic views, and BI-ready products. Validated **34/34** data-quality checks and **9/9** integration tests; delivered a six-page Streamlit dashboard.

EDUCATION & CERTIFICATES

B.Tech, Data Science — Presidency University (2019–2023) | **PGP, Data Science and Machine Learning** — Intellipaat (2024)

Certificates: PGP in Data Science and Machine Learning — Intellipaat, 24 Dec 2024 (ID 31679-462158-181209); Linux Certification Training — Intellipaat, 20 Dec 2024 (ID 31679-5930-181209); Linux Training — Intellipaat, 20 Dec 2024 (ID 31679-912747-181209); Data Science Course — Intellipaat, 4 Jan 2023 (ID 31679-8958843-181209).

Certificate IDs transcribed from supplied documents; public verification URLs were not shown on the certificates.